SAI Readiness Self-Assessment (Public)

Purpose: Help executives gauge fit for SRH reserved-capacity discovery and access fee. Score each question 0–2 (0 = Not true/Unknown, 1 = Partially true, 2 = Clearly true). A score of 14+ suggests strong fit.

1) Trust Bubble Visibility

We can quantify the gap between perceived assistant quality and ground-truth performance (e.g., conflation of fluency with correctness).

2) Failure Mode Clarity

We can name and trace our critical failure modes (e.g., sycophancy/approval-seeking, contradiction under pressure, hallucinated assurance).

3) KPI Ownership

We have internal capability to compute KPIs like ASR, HUCL, CUP, UCE, and Task Success without external content sharing.

4) Evaluation Maturity

We can run fixed-seed batches, online shadow/canary, and human adjudication on powered subsets.

5) Latency and Safety Governance

We can enforce P95 latency budgets and non-regression safety gates during controlled rollouts.

6) Data Residency and VPC Controls

We can run integrations fully inside our VPC with redaction-by-default and telemetry minimization.

7) Executive Sponsor and Decision Gate

We have an exec sponsor empowered to approve a 6–12 week pilot with a clear go/no-go KPI gate.

8) Incident Economics

We can quantify the economic upside of reducing escalations, rework, and incident risk attributable to honesty/coherence improvements.

9) Competitive Pressure

We see a clear case for exclusivity (segment/region/named competitor) tied to **first-mover advantage** on trust and alignment.

10) Capacity Alignment

Our transformation timeline can align to SRH's reserved capacity (3–5 concurrent enterprise deployments).

Scoring & Guidance

0–8: Explore readiness via internal prep; request a non-committal briefing.

9–13: Likely candidates; consider a scoped readiness baseline.

14–20: Strong fit for reserved-capacity discovery and access fee.

SRH DAO

Next step if you're a fit

SRH's SAI optimization operates through reserved capacity engagements. Given the specialized nature of distribution gap navigation and our commitment to no more than 3–5 concurrent enterprise deployments, we require a capacity reservation fee of $250,000–$500,000 to begin discovery. This fee credits toward your program investment and ensures dedicated specialist availability for your transformation timeline.

Notes

SRH does not provide dashboards or own your metrics. You compute KPIs locally using SRH event markers. The investment scales with your AI footprint and desired exclusivity. We'll design a program structure after understanding your trust bubble exposure and competitive positioning needs.

This assessment is non-binding and non-commercial; it is designed to reduce risk and accelerate an effective discovery.